

Ganesh Gupta

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

EXPERIENCE

Oracle | *Software Engineer, AI Platform*

Jun 2025 – Present

- Built and operate a multi-agent coding/tool-use platform via LangGraph, CrewAI, AutoGen, and the OpenAI Agents SDK, with MCP tool-calling and stateful agent memory via Redis checkpointing at 99%+ uptime.
- Designed and optimized agentic AI workflows across multiple LLMs (GPT, Claude, Llama, Gemini), selecting and tuning models per use case for production quality and cost.
- Provisioned and operated multi-cloud GPU clusters (AWS, GCP, CoreWeave) with NVIDIA/CUDA hardware and InfiniBand/RoCE networking for large-scale training and inference.
- Built a deployment-gated evaluation framework enforcing guardrails and quality/regression checks across agent and model outputs, blocking 15% of failed releases before production.

QureAI | *ML Infrastructure Engineer*

Mar 2026 – Present

- Extended and optimized inference serving internals (vLLM-class) with kernel-level development (CUDA/Triton), distributed KV cache, and speculative decoding to improve LLM serving performance.
- Developed multi-modal computer vision models (PyTorch), running model training and fine-tuning on multi-GPU clusters with compression and quantization for efficient inference.
- Built evaluation harnesses measuring LLM/agent output quality, drift, and regression across versions, catching failures before they reached production.

Nomura Research Institute Financial Technologies | *Software Engineer*

Sep 2019 – Aug 2023

- Built and operated 25+ production C++ and SQL-backed distributed systems for 1M+ users, sustaining 99.9% uptime over 4 years on AWS (EC2, RDS, SQS).

PROJECTS

RootCause AI

github.com/ganeshhgupta/rootcause-ai

- LangGraph multi-agent system diagnosing production incidents via tool-use chains; routes alerts to specialized sub-agents and synthesizes evidence-backed root cause reports.

Watch-your-LLM

github.com/ganeshhgupta/watch-your-llm

- Production LLM/agent eval platform with hallucination detection, drift tracking, rubric-based scoring, and regression alerting across model versions; FastAPI backend, PostgreSQL.

Distributed KV Store / VPC Control Plane Simulator

github.com/ganeshhgupta/distributed-kv-store

- Raft-consensus distributed storage engine in C++ handling concurrent writes with linearizable reads and log compaction; companion VPC control plane simulator for subnet routing.

EDUCATION

University of Texas at Arlington | *M.S. Computer Science, GPA 4.0/4.0*

Aug 2023 – May 2025

MCKV Institute of Engineering | *B.Tech Computer Science*

Jul 2016 – Jul 2020

SKILLS

Languages: Python, C++, TypeScript, Java, SQL

Agentic AI: LangGraph, CrewAI, AutoGen, OpenAI Agents SDK, MCP, Coding Agents, Multi-Agent Orchestration, Tool-Use

GPU & Inference: NVIDIA/CUDA, CUDA/Triton Kernel Development, vLLM, Distributed KV Cache, Speculative Decoding, Quantization

ML Infra: PyTorch, Model Training/Fine-Tuning (Multi-GPU), RAG, Evaluation & Regression Frameworks, Observability

Cloud & DevOps: AWS, GCP, Multi-Cloud GPU Clusters (CoreWeave), InfiniBand/RoCE, Docker, Kubernetes, CI/CD

CERTIFICATIONS

AWS Certified Developer Associate | Google Cloud Professional Machine Learning Engineer