

Ganesh Gupta

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

EXPERIENCE

Oracle | *Software Engineer, AI Platform*

Jun 2025 – Present

- Provisioned and operated multi-cloud GPU clusters (AWS, GCP, CoreWeave) with NVIDIA/CUDA hardware and InfiniBand/RoCE networking, bringing new providers and GPU hardware into production.
- Wrote infrastructure/systems-level code in Go, Rust, and C++ for scheduling, resource allocation, and fault-tolerant placement logic under GPU-scarcity constraints.
- Built observability and autoscaling for vLLM inference serving via Prometheus/Grafana; ran distributed training workloads via Slurm-based HPC scheduling, keeping fleet utilization up through node loss and provider hiccups.
- Cut p99 latency 420ms to 140ms across 50K+ daily requests on Kubernetes-managed autoscaling, optimizing distributed service throughput and cost efficiency.

QureAI | *ML Infrastructure Engineer*

Mar 2026 – Present

- Extended and optimized inference framework internals (vLLM-class serving) with kernel-level development (CUDA/Triton), distributed KV cache, and speculative decoding to improve serving performance.
- Built a model observability framework tracking drift, latency, and eval regressions in production, enabling governance-aligned rollback workflows.

Nomura Research Institute Financial Technologies | *Software Engineer*

Sep 2019 – Aug 2023

- Built and operated 25+ production C++ and SQL-backed distributed systems for 1M+ users, sustaining 99.9% uptime over 4 years on AWS (EC2, RDS, SQS).
- Standardized failure-handling contracts across 25+ services, reducing production escalations 60% in one quarter via consistent retry, timeout, and propagation semantics.

PROJECTS

Distributed KV Store / VPC Control Plane Simulator

github.com/ganeshhgupta/distributed-kv-store

- Raft-consensus distributed storage engine handling concurrent writes with linearizable reads and log compaction; companion VPC control plane simulator modeling subnet routing and resource placement.

RootCause AI

github.com/ganeshhgupta/rootcause-ai

- Multi-agent system automating production incident diagnosis via tool-use chains; routes alerts to specialized sub-agents and synthesizes root cause reports, turning operational failures into repeatable diagnostics.

Watch-your-LLM

github.com/ganeshhgupta/watch-your-llm

- Production eval and observability platform with drift detection and regression alerting across versions; FastAPI backend, PostgreSQL, deployed on Kubernetes.

EDUCATION

University of Texas at Arlington | *M.S. Computer Science, GPA 4.0/4.0*

Aug 2023 – May 2025

MCKV Institute of Engineering | *B.Tech Computer Science*

Jul 2016 – Jul 2020

SKILLS

Languages: Go, Rust, C++, Python, SQL

Infrastructure: Kubernetes, Multi-Cloud GPU Clusters (AWS, GCP, CoreWeave), InfiniBand/RoCE, Slurm HPC Scheduling, Fault-Tolerant Placement

Release & Observability: CI/CD, Prometheus/Grafana, Autoscaling, Distributed Systems Debugging

GPU Infra: NVIDIA/CUDA Hardware, vLLM, CUDA/Triton Kernel Development, Fleet Utilization Under Node Loss

CERTIFICATIONS

AWS Certified Developer Associate | Google Cloud Professional Machine Learning Engineer