

Ganesh Gupta

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

EXPERIENCE

Oracle | *Software Engineer, AI Platform*

Jun 2025 – Present

- Built large-scale **observability and telemetry** with **Prometheus/Grafana**, real-time monitoring, logging, and alerting for **vLLM** inference on multi-cloud **GPU** fleets.
- Owned the full service lifecycle from design through deployment and operation, with **on-call**, incident response, and **blameless postmortems**.
- Developed infrastructure services in **Go**, **C++**, and **Python** for **GPU** scheduling, resource allocation, and fault-tolerant placement across **AWS**, **GCP**, and **CoreWeave** fleets.
- Built **Python/FastAPI** backend services and a pluggable **distributed executor** routing workloads across **Docker** and **AWS Lambda** with an **S3**-backed artifact store.

QureAI | *ML Infrastructure Engineer*

Mar 2026 – Present

- Built and operated **Kubernetes**-based **GPU** infrastructure across clouds, writing **Go** and **Python** tooling for scheduling and fault-tolerant placement.
- Owned **ML infrastructure** end to end, from model deployment to evaluation and debugging, cutting **p99 latency 420ms to 140ms** across **50K+** daily requests.

Nomura Research Institute Financial Technologies | *Software Engineer*

Sep 2019 – Aug 2023

- Built and operated **25+** production microservices for **1M+** users on **AWS** (EC2, RDS, SQS), sustaining **99.9% uptime** over four years.
- Standardized failure-handling contracts across services and drove root-cause analysis, cutting production escalations **60%** in one quarter.
- **Profiled compute and memory bottlenecks** across **Python/C++** services and optimized batching and **concurrency**, raising throughput **2.4x** and cutting memory **27%**.

PROJECTS

Watch-your-LLM

github.com/ganeshhgupta/watch-your-llm

- **LLM evaluation** platform with rubric-based scoring, drift detection, and latency tracking across model versions; **FastAPI**, **PostgreSQL**, regression alerting.

Distributed KV Store

github.com/ganeshhgupta/distributed-kv-store

- **Raft-consensus** distributed key-value store with **linearizable reads**, leader election, and **log compaction**.

RootCause AI

github.com/ganeshhgupta/rootcause-ai

- **Multi-agent** incident-diagnosis system with **tool-use** chains that collects evidence from logs and metrics and synthesizes root-cause reports.

EDUCATION

University of Texas at Arlington | *M.S. Computer Science, GPA 4.0/4.0*

Aug 2023 – May 2025

MCKV Institute of Engineering | *B.Tech Computer Science*

Jul 2016 – Jul 2020

SKILLS

Languages: Python, Go, C++, SQL

Observability: Prometheus, Grafana, Logging, Alerting, Telemetry Pipelines, SLIs/SLOs, Capacity Management

Infrastructure: Kubernetes, Linux, Automation, Terraform, Distributed Systems

Cloud: AWS, GCP, CoreWeave, Docker

Fundamentals: Data Structures & Algorithms, Operating Systems, Database Systems, Computer Networks, HTML, CSS

Certifications: AWS Certified Developer Associate, Google Cloud Professional Machine Learning Engineer