

GANESH GUPTA

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

Experience

Oracle

Jun 2025 – Present

Software Engineer, AI Platform

- Provisioned and operated multi-cloud GPU clusters (AWS, GCP, CoreWeave) with NVIDIA/CUDA hardware, benchmarking and tuning placement for PyTorch-based training and high-availability inference on shared fleet capacity.
- Owned end-to-end design and productization of a deep-learning-backed multi-agent system, taking it from architecture through a deployment-gated evaluation pipeline enforcing automated quality and drift checks in production.
- Built observability and autoscaling for GPU inference serving using Prometheus and Grafana, keeping production ML systems reliable and within latency/safety requirements at scale.

QureAI (Healthcare AI Startup)

Jan 2024 – May 2025

ML Infrastructure Engineer

- Developed PyTorch-based computer vision deep learning models for medical imaging analysis, applying data-driven development practices to improve model accuracy and robustness on real-world clinical data.
- Led migration of model serving to async vLLM with autoscaling, cutting p99 inference latency 420ms to 140ms across 50K+ daily requests through systematic profiling and productization.

Nomura Research Institute Financial Technologies

Aug 2020 – Aug 2023

Software Engineer

- Architected and operated 25+ production Java microservices processing regulated financial transactions for 1M+ users, sustaining 99.9% uptime over 3 years on AWS (EC2, RDS, SQS).
- Drove a redesign of nightly financial data aggregation, cutting processing from 8 minutes to 90 seconds via indexed views and batch partitioning, in a domain requiring auditable, traceable transaction records.
- Led standardization of failure-handling contracts across 25+ services in partnership with cross-functional teams through code reviews and design write-ups, reducing production escalations 60% in one quarter.

Projects

Watch-your-LLM | github.com/ganeshhgupta/watch-your-llm

- Production LLM eval platform with hallucination detection, drift tracking, rubric-based scoring, and regression alerting across model versions; FastAPI backend, PostgreSQL, built to keep agentic systems trustworthy at volume.

RootCause AI | github.com/ganeshhgupta/rootcause-ai

- LangGraph multi-agent system diagnosing production incidents via tool-use chains; routes alerts to specialized sub-agents and synthesizes evidence-backed, traceable root cause reports.

Education

University of Texas at Arlington

Aug 2023 – May 2025

MS, Computer Science, GPA 4.0/4.0

MCKV Institute of Engineering

Jul 2016 – Jul 2020

BTech, Computer Science

Skills

Languages: Python, C++, Java, SQL

ML & Computer Vision: PyTorch, Computer Vision, Deep Learning Model Development, Medical Imaging Analysis, Data-Driven Development, Model Evaluation

GPU & Production ML: NVIDIA/CUDA, Multi-Cloud GPU Orchestration (AWS, GCP, CoreWeave), vLLM, Model Serving, Prometheus, Grafana

Cloud & Infra: AWS (Lambda, EC2, S3, SageMaker), GCP (Cloud Run), Docker, Kubernetes, CI/CD

Certifications

AWS Certified Developer Associate | GCP Professional MLOps Engineer