

GANESH GUPTA

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

Experience

Oracle

Jun 2025 – Present

Software Engineer, AI Platform

- Build the evaluation, guardrails, and observability layer that decides whether a release is safe to ship: a deployment-gated framework measuring quality, latency, and regression, blocking 15% of failed releases before rollout.
- Own multi-step agent workflows (LangGraph, CrewAI, AutoGen, OpenAI Agents SDK), using Codex and Claude Code daily as part of an AI-native, agentic development workflow.
- Build full-stack platform capabilities (React/TypeScript frontend, Python/FastAPI backend) and reusable, MCP-compatible SDK components used across product teams.
- Combine GraphRAG with vector-based RAG for grounded retrieval, designing REST APIs that let teams build and launch features safely end-to-end.

QureAI (AI Startup)

Mar 2026 – Present

ML Infrastructure Engineer

- Built custom evaluators tracking drift, latency, and regression across model versions, providing the trusted signal to roll changes forward or back with confidence.
- Improved semantic retrieval recall 61% to 89% on a 2K-query eval set by rebuilding the chunking pipeline with hybrid dense-sparse retrieval and re-ranking.

Nomura Research Institute Financial Technologies

Sep 2019 – Aug 2023

Software Engineer

- Architected and operated 25+ production C++ and SQL-backed microservices for 1M+ users, sustaining 99.9% uptime over 4 years on AWS (EC2, RDS, SQS).
- Drove a redesign of nightly data aggregation, cutting processing from 8 minutes to 90 seconds via indexed views and batch partitioning.
- Standardized failure-handling contracts (retry, timeout, propagation semantics) across 25+ services, reducing production escalations 60% in one quarter.
- Mentored engineers through code reviews and design write-ups, driving consistent engineering standards.

Projects

Watch-your-LLM | github.com/ganeshhgupta/watch-your-llm

- Production LLM eval platform with hallucination detection, drift tracking, rubric-based scoring, and regression alerting across model versions; FastAPI backend, PostgreSQL.

LangFetch | tinyurl.com/LangFetch

- LangGraph agent with MCP-compatible tool-use compiling natural-language questions into dependency-ordered SQL subqueries with Pinecone vector retrieval and hybrid re-ranking; deployed as a reusable MCP server for enterprise data queries.

RootCause AI | github.com/ganeshhgupta/rootcause-ai

- LangGraph multi-agent system diagnosing production incidents via tool-use chains; routes alerts to specialized sub-agents and synthesizes evidence-backed root cause reports, reducing mean time to resolution.

Education

University of Texas at Arlington

Aug 2023 – May 2025

MS, Computer Science, GPA 4.0/4.0

MCKV Institute of Engineering

Jul 2016 – Jul 2020

BTech, Computer Science

Skills

Languages: Python, TypeScript, SQL, Java, C++

Experimentation & Safety: Model/Agent Evaluation, Guardrails, Observability, Regression Tracking, Rollout Safety

Agentic AI & Dev Tools: LangGraph, CrewAI, AutoGen, OpenAI Agents SDK, Claude Code, Codex, MCP

Full-Stack & Cloud: React, FastAPI, REST APIs, PostgreSQL, Redis, AWS, GCP, Docker, CI/CD

Certifications

AWS Certified Developer Associate | GCP Professional MLOps Engineer