

GANESH GUPTA

Preferred Location: Charlotte, NC / New Jersey / Dallas, TX

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

Experience

Oracle

Jun 2025 – Present

Software Engineer, AI Platform

- Built a multi-agent AI architecture in Python and Node.js using LangGraph, the OpenAI Agents SDK, and Anthropic Claude with MCP tool-calling, orchestrating autonomous workflows for millions of daily transactions at 99%+ uptime.
- Deployed and operated AWS infrastructure (Lambda, DynamoDB, SageMaker) for real-time agentic inference at sub-200ms latency, partnering directly with product stakeholders on production integration.
- Built a deployment-gated evaluation pipeline enforcing automated guardrails around agent outputs, quality, accuracy, and drift, blocking 15% of failed releases before production launch.

QureAI (AI Startup)

Mar 2026 – Present

ML Infrastructure Engineer

- Designed and implemented a Neo4j knowledge graph modeling domain entities and relationships, deployed to production to power graph-aware retrieval and reasoning.
- Integrated AWS Bedrock-hosted foundation models, including Claude, into production agentic AI workflows for managed, scalable model access as an early engineer at this AI startup.

Nomura Research Institute Financial Technologies

Sep 2019 – Aug 2023

Software Engineer

- Architected and operated 25+ production C++ and SQL-backed microservices in distributed systems for 1M+ users, sustaining 99.9% uptime over 4 years on AWS (EC2, RDS, SQS).
- Drove a redesign of nightly data aggregation, cutting processing from 8 minutes to 90 seconds via indexed views and batch partitioning.
- Led standardization of failure-handling contracts (retry, timeout, propagation semantics) across 25+ services in partnership with cross-functional teams through code reviews, reducing production escalations 60% in one quarter.

Projects

RootCause AI | github.com/ganeshhgupta/rootcause-ai

- LangGraph multi-agent system diagnosing production incidents via tool-use chains; routes alerts to specialized sub-agents and synthesizes evidence-backed root cause reports, reducing mean time to resolution.

LangFetch | tinyurl.com/LangFetch

- LangGraph agent with Pinecone vector retrieval and hybrid re-ranking decomposes multi-table questions into dependency-ordered SQL subqueries; deployed as an MCP-compatible server for enterprise data queries.

Education

University of Texas at Arlington

Aug 2023 – May 2025

MS, Computer Science, GPA 4.0/4.0

MCKV Institute of Engineering

Jul 2016 – Jul 2020

BTech, Computer Science

Skills

Languages: Java, Python, TypeScript, SQL, C++

Agentic AI & LLM: Agentic AI, AI Architecture, Multi-Agent Systems, LangGraph, OpenAI Agents SDK, Anthropic Claude, AWS Bedrock, MCP, RAG

Full-Stack & Data: React, Node.js, FastAPI, REST APIs, GraphQL, Neo4j, DynamoDB, PostgreSQL

Cloud & MLOps: AWS (Lambda, EC2, DynamoDB, SageMaker, Bedrock), Azure, GCP (Cloud Run), Docker, Kubernetes, CI/CD

Certifications

AWS Certified Developer Associate | GCP Professional MLOps Engineer