

GANESH GUPTA

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

Experience

Oracle

Jun 2025 – Present

Software Engineer, AI Platform

- Architected high-throughput, low-latency live inference infrastructure for a multi-agent RAG system, owning it end-to-end from serving architecture to production observability at 99%+ uptime and sub-200ms latency.
- Designed and implemented observability, monitoring, and validation frameworks to detect performance gaps, latency spikes, and production drift, automatically blocking 15% of releases that failed validation before reaching production.
- Built and maintained Airflow DAGs orchestrating feature ingestion and transformation pipelines on AWS, ensuring consistency between offline and online data feeding downstream model serving.

QureAI (AI Startup)

Jan 2024 – May 2025

ML Infrastructure Engineer

- Diagnosed a live-serving retrieval-quality gap directly from real usage signals, then owned the fix end-to-end: redesigned the feature pipeline, lifting recall 61% to 89% and cutting the underlying complaint pattern from 22% to 6% of sessions.
- Led migration of synchronous model serving to async, autoscaled low-latency inference as an early engineer at this AI startup, cutting p99 latency 420ms to 140ms across 50K+ daily requests through systematic profiling and debugging.

Nomura Research Institute Financial Technologies

Aug 2020 – Aug 2023

Software Engineer

- Architected and operated 25+ production Java microservices as distributed backend infrastructure for 1M+ users, sustaining 99.9% uptime over 3 years on AWS (EC2, RDS, SQS).
- Drove a redesign of a nightly data pipeline aggregating and transforming transaction data at scale, cutting processing from 8 minutes to 90 seconds via indexed views and batch partitioning.
- Led standardization of failure-handling contracts across 25+ services in partnership with cross-functional teams through code reviews and design write-ups, reducing production escalations 60% in one quarter.

Projects

Watch-your-LLM | github.com/ganeshhgupta/watch-your-llm

- Production LLM eval platform with hallucination detection, drift tracking, rubric-based scoring, and regression alerting across model versions; FastAPI backend, PostgreSQL, built to keep agentic systems trustworthy at volume.

RootCause AI | github.com/ganeshhgupta/rootcause-ai

- LangGraph multi-agent system diagnosing production incidents via tool-use chains; routes alerts to specialized sub-agents and synthesizes evidence-backed, traceable root cause reports.

Education

University of Texas at Arlington

Aug 2023 – May 2025

MS, Computer Science, GPA 4.0/4.0

MCKV Institute of Engineering

Jul 2016 – Jul 2020

BTech, Computer Science

Skills

Languages: Python, Java, SQL, TypeScript

ML Infrastructure: Low-Latency Inference, vLLM, Feature Pipelines, Airflow, Evals, Drift/Regression Detection, Distributed Systems

Backend & APIs: FastAPI, REST APIs, PostgreSQL, RAG, Vector Databases (Pinecone, FAISS)

Cloud & Ops: AWS (Lambda, EC2, S3, RDS, SQS, SageMaker), GCP (Cloud Run), Docker, Kubernetes, CI/CD

Certifications

AWS Certified Developer Associate | GCP Professional MLOps Engineer