

Ganesh Gupta

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

EXPERIENCE

Oracle | *Software Engineer, AI Platform*

Jun 2025 – Present

- Designed and optimized agentic AI workflows across multiple LLMs (GPT, Claude, Llama, Gemini), selecting and tuning models per use case for production quality and cost.
- Built a multi-agent system via LangGraph, OpenAI Agents SDK, and MCP tool-use handling millions of daily transactions at 99%+ uptime.
- Deployed a RAG pipeline on AWS SageMaker for real-time decision support at sub-200ms p99 latency with RTO under 2 minutes.
- Built a deployment-gated evaluation and LLM Ops framework enforcing guardrails and quality checks across model outputs, blocking 15% of failed releases before production.
- Led design and code review sessions setting the architectural bar for the agentic AI codebase.

QureAI | *ML Infrastructure Engineer*

Mar 2026 – Present

- Improved RAG retrieval recall 61% to 89% on a 2K-query eval set by rebuilding the chunking pipeline with hybrid dense-sparse retrieval over a vector store (Pinecone) and re-ranking.
- Built a model observability framework tracking drift, latency, and eval regressions in production, enabling governance-aligned rollback workflows.

Nomura Research Institute Financial Technologies | *Software Engineer*

Sep 2019 – Aug 2023

- Architected 25+ production Java Spring Boot microservices for financial investment workflows serving 1M+ users at 99.9% uptime over 4 years on AWS; led architecture decisions and code-review sessions.
- Standardized failure-handling contracts across 25+ financial services, reducing production escalations 60% in one quarter via consistent retry, timeout, and propagation semantics.

PROJECTS

LangFetch

tinyurl.com/LangFetch

- LangGraph agent with MCP-compatible tool-use decomposing natural language into dependency-ordered SQL subqueries; FastAPI backend with React dashboard and PostgreSQL.

RootCause AI

github.com/ganeshhgupta/rootcause-ai

- LangGraph multi-agent system automating production incident diagnosis via tool-use chains; routes alerts to specialized sub-agents and synthesizes evidence-backed root cause reports.

Watch-your-LLM

github.com/ganeshhgupta/watch-your-llm

- Production LLM eval and observability platform with rubric-based scoring, drift detection, and regression alerting across model versions; FastAPI backend, PostgreSQL.

EDUCATION

University of Texas at Arlington | *M.S. Computer Science, GPA 4.0/4.0*

Aug 2023 – May 2025

MCKV Institute of Engineering | *B.Tech Computer Science*

Jul 2016 – Jul 2020

SKILLS

Languages: Python, Java, TypeScript, C++, SQL

AI/LLM: LLMs, Prompt Engineering, RAG, Agentic AI, LangChain, LangGraph, Vector Databases (Pinecone)

Backend & MLOps: FastAPI, Flask, REST APIs, MLOps/LLMOps, AI Observability, Kafka, Airflow

Cloud & DevOps: AWS, GCP, Docker, Kubernetes, Git, CI/CD

CERTIFICATIONS

AWS Certified Developer Associate | Google Cloud Professional Machine Learning Engineer