

GANESH GUPTA

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

Experience

Oracle

Jun 2025 – Present

Software Engineer, AI Platform

- Own production rollout of a multi-agent GenAI system using LangGraph, LangChain, and the OpenAI Agents SDK with MCP tool-calling, leading technical discussions from design through deployment.
- Deployed a RAG pipeline on AWS SageMaker for real-time decision support at sub-200ms p99 latency, partnering directly with customer-facing product stakeholders as a trusted technical advisor.
- Developed enterprise integrations and event-driven APIs connecting agentic workflows to enterprise data on AWS (Lambda, SageMaker) and Azure cloud data services.
- Designed stateful agent recovery via Redis checkpointing; interrupted workflows resume from the last committed step across distributed execution.
- Drove a deployment-gated evaluation pipeline enforcing governance and guardrails around agent access to production data, collaborating with product and data science teams to block 15% of failed releases.

QureAI (AI Startup)

Mar 2026 – Present

ML Infrastructure Engineer

- Built PySpark data processing jobs on Apache Spark for large-scale feature and dataset preparation, feeding downstream model training and evaluation pipelines.
- Improved RAG retrieval recall 61% to 89% on a 2K-query eval set by rebuilding the chunking pipeline with hybrid dense-sparse retrieval and re-ranking.

Nomura Research Institute Financial Technologies

Sep 2019 – Aug 2023

Software Engineer

- Architected and operated 25+ production C++ and SQL-backed microservices in distributed systems for 1M+ users, sustaining 99.9% uptime over 4 years on AWS (EC2, RDS, SQS).
- Drove a redesign of nightly data aggregation, cutting processing from 8 minutes to 90 seconds via indexed views and batch partitioning.
- Standardized failure-handling contracts (retry, timeout, propagation semantics) across 25+ services, reducing production escalations 60% in one quarter.
- Mentored engineers through code reviews and design write-ups, driving consistent engineering standards across the team.

Projects

LangFetch | tinyurl.com/LangFetch

- LangGraph agent compiling natural-language questions into dependency-ordered SQL subqueries with Pinecone vector retrieval and hybrid re-ranking; deployed as an MCP-compatible server for enterprise data queries.

RootCause AI | github.com/ganeshhgupta/rootcause-ai

- LangGraph multi-agent system diagnosing production incidents via tool-use chains; routes alerts to specialized sub-agents and synthesizes evidence-backed root cause reports, reducing mean time to resolution.

Watch-your-LLM | github.com/ganeshhgupta/watch-your-llm

- Production LLM eval platform with hallucination detection, drift tracking, rubric-based scoring, and regression alerting across model versions; FastAPI backend, PostgreSQL.

Education

University of Texas at Arlington

Aug 2023 – May 2025

MS, Computer Science, GPA 4.0/4.0

MCKV Institute of Engineering

Jul 2016 – Jul 2020

BTech, Computer Science

Skills

Languages: Python, SQL, TypeScript, Java, C++

Data Engineering: PySpark, Apache Spark, Airflow, ETL/ELT, Data Pipelines, pandas, scikit-learn

GenAI & Agentic AI: Agentic AI, RAG, Text2SQL, LangChain, LangGraph, DSPy, Hugging Face, PyTorch, MCP

Cloud & MLOps: AWS (Lambda, SageMaker), Azure, GCP, Docker, Kubernetes, Git, CI/CD

Certifications

AWS Certified Developer Associate | GCP Professional MLOps Engineer