

Ganesh Gupta

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

EXPERIENCE

Oracle | *Software Engineer, AI Platform*

Jun 2025 – Present

- Designed and optimized agentic AI workflows across multiple foundation models (GPT, Claude, Llama, Gemini), studying model-internal behavior differences to select and tune models per use case.
- Built a multi-agent system via OpenAI Agents SDK and MCP tool-use handling millions of daily transactions at 99%+ uptime, prototyping and iterating on agent tool-use and failure modes.
- Deployed a RAG pipeline on AWS SageMaker for real-time decision support at sub-200ms p99 latency, running controlled experiments on retrieval and generation tradeoffs.
- Participated in design and code reviews enforcing rigor and testability standards across the agentic AI codebase.

QureAI | *ML Infrastructure Engineer*

Mar 2026 – Present

- Built a model observability framework tracking drift, latency, and eval regressions in production, designing rigorous evaluation methodology that gates governance-aligned rollback decisions.
- Improved RAG retrieval recall 61% to 89% on a 2K-query eval set through systematic experimentation with hybrid dense-sparse retrieval and re-ranking.

Nomura Research Institute Financial Technologies | *Software Engineer*

Sep 2019 – Aug 2023

- Built and operated 25+ production microservices in Java, C++, and SQL-backed distributed systems for 1M+ users, sustaining 99.9% uptime over 3 years on AWS (EC2, RDS, SQS).
- Standardized failure-handling contracts across 25+ financial services, reducing production escalations 60% in one quarter via consistent retry, timeout, and propagation semantics.

PROJECTS

RootCause AI

github.com/ganeshhgupta/rootcause-ai

- LangGraph multi-agent system automating production incident diagnosis via tool-use chains; routes findings to specialized sub-agents and synthesizes evidence-backed root cause reports.

Watch-your-LLM

github.com/ganeshhgupta/watch-your-llm

- Built a rigorous LLM evaluation harness with rubric-based scoring, drift detection, and regression tracking across model versions for production governance.

LLaMA 2 Fine-Tuning Pipeline for Financial Documents

github.com/ganeshhgupta/llama2-finance-finetuning

- Post-trained LLaMA 2 internals with LoRA and PEFT for domain-specific document understanding, running systematic hyperparameter experiments to study capacity-performance tradeoffs.

EDUCATION

University of Texas at Arlington | *M.S. Computer Science, GPA 4.0/4.0*

Aug 2023 – May 2025

MCKV Institute of Engineering | *B.Tech Computer Science*

Jul 2016 – Jul 2020

SKILLS

Languages: Python, TypeScript, Java, C++, SQL

AI/Research: LLMs & Foundation Models (GPT, Claude, Llama, Gemini), Agentic AI, Multi-Agent Systems, LoRA/PEFT Fine-Tuning, Model Evaluation & Rigorous Eval Harnesses, RAG

Tooling: LangChain, LangGraph, OpenAI Agents SDK, MCP, PostgreSQL

Cloud & DevOps: AWS (EC2, S3, SageMaker), GCP (Cloud Run), Docker, Kubernetes, CI/CD

CERTIFICATIONS

AWS Certified Developer Associate | Google Cloud Professional Machine Learning Engineer