

Ganesh Gupta

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

EXPERIENCE

Oracle | *Software Engineer, AI Platform*

Jun 2025 – Present

- Built observability and autoscaling for **vLLM** inference serving with **Prometheus/Grafana**, and trained and fine-tuned **PyTorch** models on multi-GPU clusters.
- Provisioned and operated multi-cloud **GPU clusters** (AWS, GCP, CoreWeave) with **NVIDIA/CUDA** and **InfiniBand**, running distributed training on **Slurm** and **Kubernetes**.
- Ran **A/B tests** on agent and model changes in production before full rollout, pairing them with **deployment-gated evals** to decide releases on measured impact.
- Developed and deployed **deep learning models** in **PyTorch** and **TensorFlow** with automated training workflows and **MLOps CI/CD** on **AWS SageMaker**.

QureAI | *ML Infrastructure Engineer*

Mar 2026 – Present

- Built **PySpark** jobs for large-scale dataset preparation feeding model training and **evaluation pipelines** in a regulated **healthcare** product.
- Cut **p99 latency from 420ms to 140ms** across **50K+ daily requests** by moving LLM serving to an **async, autoscaled** architecture.
- Built a model **observability** platform tracking drift, latency, and eval regressions, with **automated CI/CD rollback** that cut manual incident toil.

Nomura Research Institute Financial Technologies | *Software Engineer*

Sep 2019 – Aug 2023

- Integrated **ML models (scikit-learn, TensorFlow)** into production workflows for **1M+ users** at **99.9% uptime**, with data pipelines processing **50K+ records** on AWS.
- Architected and operated **25+ production Python and Java microservices** for **1M+ users** in **financial services** at **99.9% uptime** on **AWS**.
- Cut nightly batch aggregation from **8 minutes to 90 seconds** by optimizing **SQL and PL/SQL** on a financial RDBMS with indexed views and partitioning.
- Cut production escalations **60%** in one quarter by standardizing API contracts and failure handling, driving **root cause analysis** and post-incident reviews.

PROJECTS

Watch-your-LLM

github.com/ganeshhgupta/watch-your-llm

- **LLM evaluation** platform with rubric-based scoring, **hallucination detection**, regression alerting, a **FastAPI** backend, and a **React/TypeScript** dashboard.

Distributed KV Store

github.com/ganeshhgupta/distributed-kv-store

- **Raft-consensus** key-value store with **linearizable reads** and **log compaction**, plus a companion VPC control plane simulator.

RootCause AI

github.com/ganeshhgupta/rootcause-ai

- **LangGraph** multi-agent system that routes production alerts to specialized sub-agents via **tool-use** chains and writes structured root-cause reports, cutting **MTTR 65%**.

EDUCATION

University of Texas at Arlington | *M.S. Computer Science, GPA 4.0/4.0*

Aug 2023 – May 2025

MCKV Institute of Engineering | *B.Tech Computer Science*

Aug 2016 – Jul 2020

SKILLS

Languages: Python, Java, SQL, C++, Go

ML: PyTorch, TensorFlow, scikit-learn, Feature Engineering, Model Training, Online Serving, A/B Testing, Model Monitoring

Data & Distributed: Spark/PySpark, Kafka, Airflow, NoSQL, Redis, DynamoDB, PostgreSQL

Infra: Kubernetes, Slurm, vLLM, NVIDIA GPUs, AWS, GCP, Prometheus/Grafana, Claude Code, Codex

Certifications: AWS Certified Developer Associate, GCP Professional MLOps Engineer