

Ganesh Gupta

817-822-1131 | ganeshg12ta@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

EXPERIENCE

Oracle | *Software Engineer, AI Platform*

Jun 2025 – Present

- Wrote production **agentic code** orchestrating **LLMs** and tool-using agents (**LangGraph**, **OpenAI Agents SDK**, **MCP**) with **state machines**, **contextual memory**, and **Redis checkpointing** for resumable, auditable runs.
- Built **asynchronous workflows** and backend services in **Python/FastAPI** with a pluggable **distributed executor** across **Docker** and **AWS Lambda**, designed around failure modes and clear data provenance.
- Built **control-plane** style **observability** with **Prometheus/Grafana** and streamed live agent state to **React/TypeScript** UIs over **SSE**, so operators can supervise and intervene in agent runs.
- Shipped a deployment-gated **evaluation** framework with **guardrails** and approval gates, blocking **15%** of failed releases before production.

QureAI | *ML Infrastructure Engineer*

Mar 2026 – Present

- Owned **ML infrastructure** end to end, from model deployment to evaluation and debugging, cutting **p99 latency 420ms to 140ms** across **50K+** daily requests.
- Improved retrieval recall **61% to 89%** on a 2K-query eval set by rebuilding chunking with **hybrid dense-sparse retrieval** and re-ranking.

Nomura Research Institute Financial Technologies | *Software Engineer*

Sep 2019 – Aug 2023

- Built **React/TypeScript** front ends and **Java/Spring** backends end to end, from domain models and APIs to high-density **dashboards**, for **1M+** users.
- Built and operated **25+** production microservices for **1M+** users on **AWS** (EC2, RDS, SQS), sustaining **99.9% uptime** over four years.
- Standardized failure-handling contracts across services and drove root-cause analysis, cutting production escalations **60%** in one quarter.

PROJECTS

RootCause AI

github.com/ganeshhgupta/rootcause-ai

- **Multi-agent** incident-diagnosis system with **tool-use** chains that collects evidence from logs and metrics and synthesizes root-cause reports.

Watch-your-LLM

github.com/ganeshhgupta/watch-your-llm

- **LLM evaluation** platform with rubric-based scoring, drift detection, and latency tracking across model versions; **FastAPI**, **PostgreSQL**, regression alerting.

LangFetch

tinyurl.com/LangFetch

- **LangGraph** agent decomposing natural-language questions into dependency-ordered SQL subqueries with hybrid retrieval and re-ranking; deployable as an **MCP** server.

EDUCATION

University of Texas at Arlington | *M.S. Computer Science, GPA 4.0/4.0*

Aug 2023 – May 2025

MCKV Institute of Engineering | *B.Tech Computer Science*

Jul 2016 – Jul 2020

SKILLS

Languages: Python, TypeScript, JavaScript, Java, SQL

Agentic AI: LLM Orchestration, Tool Use, State Machines, Memory, Human-in-the-Loop, Guardrails, LangGraph, MCP

Full Stack: React, TypeScript, Data Visualization, FastAPI, Spring, Async Workflows, Event Streams (Kafka), PostgreSQL, Redis

Cloud & DevOps: AWS, GCP, Kubernetes, Docker, Terraform, CI/CD, Prometheus, Grafana

Fundamentals: Data Structures & Algorithms, Operating Systems, Database Systems, Computer Networks, HTML, CSS

Certifications: AWS Certified Developer Associate, Google Cloud Professional Machine Learning Engineer