

Ganesh Gupta

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

EXPERIENCE

Oracle | *Software Engineer, AI Platform*

Jun 2025 – Present

- Built **multi-agent system** via **OpenAI Agents SDK** + **MCP** tool-use handling **millions of daily transactions** at **99%+ uptime**.
- Designed and optimized **agentic AI** workflows across **GPT**, **Claude**, **Llama**, and **Gemini**, selecting and tuning models per use case for production **quality and cost**.
- Engineered pluggable **distributed executor** routing pipeline nodes to **Docker** or **AWS Lambda**, backed by **S3** artifact store eliminating redundant recomputation.
- Built **evaluation harnesses** measuring LLM output quality, safety, **hallucination rates**, and grounding across agent versions; iterated on **eval results** and production usage data.

QureAI | *ML Infrastructure Engineer*

Mar 2026 – Present

- Improved **RAG retrieval recall 61% to 89%** on a 2K-query eval set by rebuilding chunking with **hybrid dense-sparse retrieval** and **re-ranking**.
- Built **model observability** framework tracking **drift**, latency, and **eval regressions** in production, enabling governance-aligned **rollback** workflows.

Nomura Research Institute Financial Technologies | *Software Engineer*

Sep 2019 – Aug 2023

- Built and operated **25+ production microservices** in **C++** and **SQL**-backed **distributed systems** for **1M+ users**, sustaining **99.9% uptime** over 4 years on **AWS** (EC2, RDS, SQS).
- Standardized failure-handling contracts across **25+ financial services**, reducing production escalations **60%** in one quarter via consistent **retry**, **timeout**, and **propagation** semantics.
- Standardized **REST API** contracts across **25+ Java services** with consistent **retry**, **timeout**, and **OAuth**-aligned auth; reduced cross-service escalations **60%** within one quarter.

PROJECTS

RootCause AI

github.com/ganeshhgupta/rootcause-ai

- LangGraph** multi-agent system diagnosing **production incidents** via tool-use chains; routes alerts to specialized sub-agents and synthesizes **root cause reports**, reducing **MTTR**.

LangFetch

tinyurl.com/LangFetch

- LangGraph** agent decomposes multi-table questions into dependency-ordered **SQL subqueries** with intermediate result passing; deployable as **MCP-compatible** server.

Watch-your-LLM

github.com/ganeshhgupta/watch-your-llm

- Production **LLM eval platform** with rubric-based scoring, **drift detection**, and latency tracking across model versions; **FastAPI**, **React/TypeScript**, **PostgreSQL**.

EDUCATION

University of Texas at Arlington | *M.S. Computer Science, GPA 4.0/4.0*

Aug 2023 – May 2025

MCKV Institute of Engineering | *B.Tech Computer Science*

Jul 2016 – Jul 2020

SKILLS

Languages: Python, TypeScript, Java, SQL, HTML/CSS

Agentic AI: OpenAI Agents SDK, LangGraph, MCP, Multi-Agent Orchestration, RAG, Vector Databases

Integration: REST APIs, Webhooks, OAuth, Enterprise Integration

Infra: AWS, GCP, Docker, Kubernetes, PostgreSQL, Redis, CI/CD

CERTIFICATIONS

AWS Certified Developer Associate | GCP Professional MLOps Engineer