

GANESH GUPTA

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

Experience

Oracle

Jun 2025 – Present

Software Engineer, AI Platform

- Lead multi-agent system design (LangGraph, CrewAI, AutoGen, OpenAI Agents SDK) with advanced RAG: multi-stage retrieval, hybrid search, and re-ranking for multi-hop reasoning over enterprise data.
- Own AI evaluation and performance benchmarking (A/B testing, model optimization) alongside cost optimization across GPU-backed model serving infrastructure.
- Built AI observability covering logging, tracing, monitoring, and automated alerting, with responsible AI review and model validation gating 15% of releases before production.
- Architect cloud-native platform APIs (Python, FastAPI, React) with distributed data ingestion pipelines feeding feature stores for downstream model consumers.

QureAI (AI Startup)

Mar 2026 – Present

ML Infrastructure Engineer

- Built custom evaluators for RAG/agent workflows tracking drift, latency, and regression across model versions, catching failures before they reached production.
- Improved semantic retrieval recall 61% to 89% on a 2K-query eval set by rebuilding the chunking pipeline with hybrid dense-sparse retrieval and re-ranking.

Nomura Research Institute Financial Technologies

Sep 2019 – Aug 2023

Software Engineer

- Architected and operated 25+ production C++ and SQL-backed microservices for 1M+ users, sustaining 99.9% uptime over 4 years on AWS (EC2, RDS, SQS).
- Drove a redesign of nightly data aggregation, cutting processing from 8 minutes to 90 seconds via indexed views and batch partitioning.
- Standardized failure-handling contracts (retry, timeout, propagation semantics) across 25+ distributed services, reducing production escalations 60% in one quarter.
- Mentored engineers through code reviews and design write-ups, driving consistent engineering standards.

Projects

Watch-your-LLM | github.com/ganeshhgupta/watch-your-llm

- Production LLM eval platform with hallucination detection, drift tracking, rubric-based scoring, and regression alerting across model versions; FastAPI backend, PostgreSQL.

LangFetch | tinyurl.com/LangFetch

- Agent compiling natural-language questions into dependency-ordered SQL subqueries with vector retrieval and hybrid re-ranking; TypeScript/Node.js backend, React dashboard, PostgreSQL.

RootCause AI | github.com/ganeshhgupta/rootcause-ai

- Multi-agent system diagnosing production incidents via tool-use chains; routes alerts to specialized sub-agents and synthesizes evidence-backed root cause reports, reducing mean time to resolution.

Education

University of Texas at Arlington

Aug 2023 – May 2025

MS, Computer Science, GPA 4.0/4.0

MCKV Institute of Engineering

Jul 2016 – Jul 2020

BTech, Computer Science

Skills

Languages: Python, TypeScript, SQL, Java, C++

Agentic AI & Retrieval: LangGraph, CrewAI, AutoGen, OpenAI Agents SDK, Advanced RAG, Multi-Stage Retrieval, Re-Ranking, Vector Databases, Embedding Models

AI Platform Engineering: AI Evaluation, A/B Testing, Model Optimization, Observability, Logging/Tracing, Responsible AI, Bias Mitigation

Cloud & Infra: AWS, GCP, FastAPI, React, Docker, Kubernetes, GPU Infrastructure, CI/CD

Certifications

AWS Certified Developer Associate | GCP Professional MLOps Engineer