

GANESH GUPTA

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

Experience

Oracle

Jun 2025 – Present

Software Engineer, AI Platform

- Designed and built production agentic AI systems in Python using LangGraph and the OpenAI Agents SDK with MCP tool integration: multi-agent orchestration, memory management, and autonomous decision-making at 99%+ uptime.
- Built a deployment-gated evaluation pipeline enforcing automated guardrails around agent outputs, quality, accuracy, and drift, blocking 15% of failed releases and validating AI-generated outputs before production launch.
- Built and operated AWS serverless infrastructure (Lambda, SageMaker, DynamoDB) for real-time agentic inference at sub-200ms latency, moving fast from prototype to production alongside product stakeholders.

QureAI (AI Startup)

Mar 2026 – Present

ML Infrastructure Engineer

- Diagnosed a user-facing RAG retrieval gap directly from real usage signals as an early engineer at this AI startup, then owned the fix end-to-end in Python: redesigned chunking, embeddings, and vector search strategy via systematic model evaluation, lifting recall 61% to 89%.
- Migrated synchronous PyTorch-based model serving to async vLLM with autoscaling, cutting p99 inference latency 420ms to 140ms across 50K+ daily requests, moving a production ML system from prototype to scale.

Nomura Research Institute Financial Technologies

Sep 2019 – Aug 2023

Software Engineer

- Architected and operated 25+ production C++ and SQL-backed microservices in distributed systems for 1M+ users, sustaining 99.9% uptime over 4 years on AWS (EC2, RDS, SQS).
- Drove a redesign of nightly data aggregation, cutting processing from 8 minutes to 90 seconds via indexed views and batch partitioning.
- Led standardization of failure-handling contracts and code-review practices across 25+ services in partnership with cross-functional teams, reducing production escalations 60% in one quarter.

Projects

LangFetch | tinyurl.com/LangFetch

- LangGraph agent with Pinecone vector retrieval and hybrid re-ranking decomposes multi-table questions into dependency-ordered SQL subqueries; deployed as an MCP-compatible server for enterprise data queries.

Watch-your-LLM | github.com/ganeshhgupta/watch-your-llm

- Production LLM eval platform with hallucination detection, drift tracking, rubric-based scoring, and regression alerting across model versions; FastAPI backend, PostgreSQL.

Education

University of Texas at Arlington

Aug 2023 – May 2025

MS, Computer Science, GPA 4.0/4.0

MCKV Institute of Engineering

Jul 2016 – Jul 2020

BTech, Computer Science

Skills

Languages: Python, SQL, Java, TypeScript, C++

Agentic AI & LLM: Agentic AI, LangGraph, OpenAI Agents SDK, MCP, RAG, Vector Search & Embeddings (Pinecone, FAISS), AI/Agent Guardrails, Prompt Engineering, Model Evaluation, Multi-Agent Orchestration

Backend & Data: FastAPI, REST APIs, PostgreSQL, MongoDB, Data Pipelines, Distributed Systems

Cloud & MLOps: AWS (Lambda, EC2, S3, SageMaker, DynamoDB), GCP (Cloud Run), Docker, Kubernetes, CI/CD, Code Review

Certifications

AWS Certified Developer Associate | GCP Professional MLOps Engineer