

GANESH GUPTA

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

Experience

Oracle

Jun 2025 – Present

Software Engineer, AI Platform

- Own production-grade multi-agent orchestration (LangGraph, CrewAI, AutoGen, OpenAI Agents SDK) with MCP-compatible tool integrations and stateful agent memory via Redis checkpointing at 99%+ uptime.
- Build agentic systems across multiple AI SDKs (OpenAI, Anthropic Claude, Google Gemini/Llama), combining GraphRAG with vector-based RAG for grounded enterprise retrieval.
- Built a deployment-gated evaluation framework with AI observability, guardrails, and monitoring, blocking 15% of failed releases before production.
- Build enterprise AI platform APIs (Python, TypeScript) partnering with product and business stakeholders end-to-end.

QureAI (AI Startup)

Mar 2026 – Present

ML Infrastructure Engineer

- Built custom evaluators for RAG/agent workflows tracking drift, latency, and regression across model versions, catching failures before they reached production.
- Improved semantic retrieval recall 61% to 89% on a 2K-query eval set by rebuilding the chunking pipeline with hybrid dense-sparse retrieval and re-ranking.

Nomura Research Institute Financial Technologies

Sep 2019 – Aug 2023

Software Engineer

- Architected and operated 25+ production C++ and SQL-backed microservices for 1M+ users, sustaining 99.9% uptime over 4 years on AWS (EC2, RDS, SQS).
- Drove a redesign of nightly data aggregation, cutting processing from 8 minutes to 90 seconds via indexed views and batch partitioning.
- Standardized failure-handling contracts (retry, timeout, propagation semantics) across 25+ services, reducing production escalations 60% in one quarter.
- Mentored engineers through code reviews and design write-ups, driving consistent engineering standards.

Projects

RootCause AI | github.com/ganeshhgupta/rootcause-ai

- LangGraph multi-agent system diagnosing production incidents via tool-use chains; routes alerts to specialized sub-agents and synthesizes evidence-backed root cause reports, reducing mean time to resolution.

LangFetch | tinyurl.com/LangFetch

- LangGraph agent with MCP-compatible tool-use compiling natural-language questions into dependency-ordered SQL subqueries with Pinecone vector retrieval and hybrid re-ranking; deployed as an MCP server for enterprise data queries.

Watch-your-LLM | github.com/ganeshhgupta/watch-your-llm

- Production LLM eval platform with hallucination detection, drift tracking, rubric-based scoring, and regression alerting across model versions; FastAPI backend, PostgreSQL.

Education

University of Texas at Arlington

Aug 2023 – May 2025

MS, Computer Science, GPA 4.0/4.0

MCKV Institute of Engineering

Jul 2016 – Jul 2020

BTech, Computer Science

Skills

Languages: Python, TypeScript, SQL, Java, C++

Agentic AI & SDKs: LangGraph, LangChain, LlamaIndex, CrewAI, AutoGen, OpenAI Agents SDK, MCP, OpenAI, Claude, Gemini

Retrieval & Observability: RAG, GraphRAG, Model/Agent Evaluation, Guardrails, AI Observability

Cloud & Platform: AWS, GCP, Docker, Kubernetes, PostgreSQL, Redis, CI/CD

Certifications

AWS Certified Developer Associate | GCP Professional MLOps Engineer