

# GANESH GUPTA

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

## Experience

### Oracle

Jun 2025 – Present

*Software Engineer, AI Platform*

- Serve as technical anchor for the agentic AI platform, defining architecture for multi-agent and sub-agent orchestration using LangGraph, LangChain, and the OpenAI Agents SDK with MCP tool integration at 99%+ uptime.
- Built a deployment-gated evaluation pipeline enforcing guardrails, AI governance, and observability around agent outputs, personally building the most technically challenging components to block 15% of failed releases.
- Developed full-stack enterprise integrations (React frontend, Node.js/Python backend) and event-driven APIs on multi-tenant AWS infrastructure (Lambda, SageMaker) and Azure at sub-200ms latency.
- Designed stateful agent recovery and memory systems via Redis checkpointing; interrupted workflows resume from the last committed step across distributed execution.
- Established engineering standards for CI/CD, testing, and security across the agentic platform, mentoring engineers through architecture reviews and design write-ups.

### QureAI (AI Startup) (Intern)

Mar 2026 – Present

*ML Infrastructure Engineer*

- Integrated AWS Bedrock-hosted foundation models, including Claude, into production agentic AI workflows for managed, scalable model access as an early engineer at this AI startup.
- Improved RAG retrieval recall 61% to 89% on a 2K-query eval set by rebuilding the chunking pipeline with hybrid dense-sparse retrieval and re-ranking.

### Nomura Research Institute Financial Technologies

Sep 2019 – Aug 2023

*Software Engineer*

- Architected and operated 25+ production C++ and SQL-backed microservices in distributed systems for 1M+ users, sustaining 99.9% uptime over 4 years on AWS (EC2, RDS, SQS).
- Drove a redesign of nightly data aggregation, cutting processing from 8 minutes to 90 seconds via indexed views and batch partitioning.
- Standardized failure-handling contracts (retry, timeout, propagation semantics) across 25+ services, reducing production escalations 60% in one quarter.
- Mentored engineers through code reviews and design write-ups, driving consistent engineering standards across the team.

## Projects

### LangFetch | [tinyurl.com/LangFetch](https://tinyurl.com/LangFetch)

- LangGraph agent compiling natural-language questions into dependency-ordered SQL subqueries with Pinecone vector retrieval and hybrid re-ranking; deployed as an MCP-compatible server for enterprise data queries.

### RootCause AI | [github.com/ganeshhgupta/rootcause-ai](https://github.com/ganeshhgupta/rootcause-ai)

- LangGraph multi-agent system diagnosing production incidents via tool-use chains; routes alerts to specialized sub-agents and synthesizes evidence-backed root cause reports, reducing mean time to resolution.

### Watch-your-LLM | [github.com/ganeshhgupta/watch-your-llm](https://github.com/ganeshhgupta/watch-your-llm)

- Production LLM eval platform with hallucination detection, drift tracking, rubric-based scoring, and regression alerting across model versions; FastAPI backend, PostgreSQL.

## Education

### University of Texas at Arlington

Aug 2023 – May 2025

*MS, Computer Science, GPA 4.0/4.0*

### MCKV Institute of Engineering

Jul 2016 – Jul 2020

*BTech, Computer Science*

## Skills

**Languages:** Python, Java, TypeScript, SQL, C++

**Agentic AI & LLM:** Agentic AI, Multi-Agent & Sub-Agent Orchestration, LangGraph, LangChain, MCP, RAG, Memory Systems, AI Governance

**Platform & Security:** React, Node.js, FastAPI, REST APIs, Microservices, RBAC, Guardrails

**Cloud & Infra:** AWS, GCP (Vertex AI, BigQuery), Azure, Docker, Kubernetes, Terraform, CI/CD

## Certifications

AWS Certified Developer Associate | GCP Professional MLOps Engineer