

Ganesh Gupta

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

EXPERIENCE

Oracle | *Software Engineer, AI Platform*

Jun 2025 – Present

- Built **multi-agent system** via **OpenAI Agents SDK** + **MCP** handling **millions of daily transactions** at **99%+ uptime**.
- Orchestrated **multi-agent** workflows with **CrewAI**, **AutoGen**, **LangGraph**, **OpenAI Agents SDK**; compared **open-source agentic frameworks**.
- Designed stateful recovery via per-node **Redis checkpointing**; interrupted workflows **resume from the last committed step** across **distributed execution**.
- Engineered pluggable **distributed executor** routing nodes to **Docker** or **AWS Lambda**, backed by **S3** artifact store eliminating recomputation.
- Tuned **agentic AI** workflows across **GPT**, **Claude**, **Llama**, **Gemini**; used **Codex/Claude Code** for **AI-assisted engineering**.

QureAI | *ML Infrastructure Engineer*

Mar 2026 – Present

- Built **model observability** tracking **drift**, latency, and **eval regressions** in production, enabling governance-aligned **rollback**.
- Integrated **Bedrock Guardrails content-filtering** and **safety controls** into production agentic workflows on clinical data.
- Cut **p99 latency 420ms to 140ms** across **50K+ daily requests** via async **vLLM** on **GCP Cloud Run**.

Nomura Research Institute Financial Technologies | *Software Engineer*

Sep 2019 – Aug 2023

- Built and operated **25+ production microservices** in **C++** and **SQL-backed distributed systems** for **1M+ users** at **99.9% uptime** on **AWS**.
- Cut nightly financial aggregation blocking **settlement** from **8 minutes to 90 seconds** via **indexed views** and **batch partitioning**.
- Standardized failure-handling across **25+ financial services**, cutting escalations **60%** via consistent **retry**, **timeout**, and **propagation** semantics.

PROJECTS

RootCause AI

github.com/ganeshhgupta/rootcause-ai

- LangGraph** multi-agent system diagnosing **production incidents**; routes alerts to specialized **sub-agents**, synthesizes **root cause reports**.

LangFetch

tinyurl.com/LangFetch

- LangGraph** agent decomposing questions into dependency-ordered **SQL subqueries**; deployable as **MCP-compatible server**.

Watch-your-LLM

github.com/ganeshhgupta/watch-your-llm

- Production **LLM eval platform** with rubric scoring, **drift detection**, latency tracking; **FastAPI**, **React/TypeScript**, **PostgreSQL**.

EDUCATION

University of Texas at Arlington | *M.S. Computer Science, GPA 4.0/4.0*

Aug 2023 – May 2025

MCKV Institute of Engineering | *B.Tech Computer Science*

Jul 2016 – Jul 2020

SKILLS

Languages & Fundamentals: Python, C++, Java, Go, SQL, TypeScript, HTML/CSS, Distributed Systems, DS&A

Agentic AI: LangGraph, LangChain, OpenAI Agents SDK, CrewAI, AutoGen, MCP, Claude Code, Codex, RAG, LLM Memory, Evals

ML Infra: vLLM, TensorRT-LLM, PyTorch, Redis, PostgreSQL, Airflow, Prometheus, Grafana, REST APIs

Cloud & Infra: AWS (EC2, RDS, S3, Lambda, SQS), GCP, Azure, Docker, Kubernetes, GitHub Actions

CERTIFICATIONS

AWS Certified Developer Associate | GCP Professional MLOps Engineer