

GANESH GUPTA

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

Experience

Oracle

Jun 2025 – Present

Software Engineer, AI Platform

- Lead architecture and solution design for multi-agent AI systems using LangGraph, LangChain, and the OpenAI Agents SDK with MCP tool-calling, handling millions of daily transactions at 99%+ uptime.
- Rapidly prototype and ship PoCs for new GenAI and agentic AI use cases across multiple LLMs (GPT, Claude, Llama, Gemini), moving fast from architecture through production.
- Developed full-stack enterprise integrations (React frontend, Node.js/Python backend) and event-driven APIs connecting agentic workflows to AWS (Lambda, SageMaker) and Azure at sub-200ms latency.
- Designed stateful agent recovery via Redis checkpointing; interrupted workflows resume from the last committed step across distributed execution.
- Drove a deployment-gated evaluation pipeline governing agent output quality with platform and security teams, blocking 15% of failed releases.

QureAI (AI Startup)

Mar 2026 – Present

ML Infrastructure Engineer

- Integrated AWS Bedrock-hosted foundation models, including Claude, into production agentic AI workflows for managed, scalable model access as an early engineer at this AI startup.
- Improved RAG retrieval recall 61% to 89% on a 2K-query eval set by rebuilding the chunking pipeline with hybrid dense-sparse retrieval and re-ranking.

Nomura Research Institute Financial Technologies

Sep 2019 – Aug 2023

Software Engineer

- Architected and operated 25+ production C++ and SQL-backed microservices in distributed systems for 1M+ users, sustaining 99.9% uptime over 4 years on AWS (EC2, RDS, SQS).
- Drove a redesign of nightly data aggregation, cutting processing from 8 minutes to 90 seconds via indexed views and batch partitioning.
- Standardized failure-handling contracts (retry, timeout, propagation semantics) across 25+ services, reducing production escalations 60% in one quarter.
- Mentored engineers through code reviews and design write-ups, driving consistent engineering standards across the team.

Projects

LangFetch | tinyurl.com/LangFetch

- LangGraph agent compiling natural-language questions into dependency-ordered SQL subqueries with Pinecone vector retrieval and hybrid re-ranking; deployed as an MCP-compatible server for enterprise data queries.

RootCause AI | github.com/ganeshhgupta/rootcause-ai

- LangGraph multi-agent system diagnosing production incidents via tool-use chains; routes alerts to specialized sub-agents and synthesizes evidence-backed root cause reports, reducing mean time to resolution.

Watch-your-LLM | github.com/ganeshhgupta/watch-your-llm

- Production LLM eval platform with hallucination detection, drift tracking, rubric-based scoring, and regression alerting across model versions; FastAPI backend, PostgreSQL.

Education

University of Texas at Arlington

Aug 2023 – May 2025

MS, Computer Science, GPA 4.0/4.0

MCKV Institute of Engineering

Jul 2016 – Jul 2020

BTech, Computer Science

Skills

Languages: Python, Java, TypeScript, SQL, C++

Agentic AI & LLM: Agentic AI, Multi-Agent Architectures, LangGraph, LangChain, GPT, Claude, Llama, Gemini, MCP, RAG

Full-Stack & Integration: React, Node.js, FastAPI, REST APIs, Event-Driven Architecture, Microservices

Cloud & Infra: AWS (Lambda, EC2, RDS, SQS, SageMaker, Bedrock), Azure, Docker, Kubernetes, CI/CD

Certifications

AWS Certified Developer Associate | GCP Professional MLOps Engineer