

# Ganesh Gupta

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

## EXPERIENCE

---

### Oracle | *Software Engineer, AI Platform*

Jun 2025 – Present

- Owned frontend-to-backend delivery of AI agent workflows using **React**, **TypeScript**, and **Node.js**; built a real-time **SSE** UI serving **millions of daily transactions** at **99%+ uptime**.
- Replaced polling with **SSE** streaming for live dashboards, cutting redundant **RESTful API** calls **90%** and resolving production issues before user impact.
- Engineered executor routing pipeline nodes to **Docker** and **AWS Lambda** via typed **REST APIs**, backed by an **S3** artifact store at **sub-200ms p99**.
- Designed stateful recovery via per-node **Redis checkpointing**, so interrupted workflows **resume from the last committed step**; shipped via **GitHub Actions CI/CD** with **code reviews**.

### QureAI | *ML Infrastructure Engineer*

Mar 2026 – Present

- Cut **p99 inference latency 420ms to 140ms** across **50K+ daily requests** on **GCP Cloud Run** with **Kubernetes** autoscaling.
- Containerized services with **Docker** and built production observability with automated **CI/CD rollback** and **failure alerting**.

### Nomura Research Institute Financial Technologies | *Software Engineer*

Sep 2019 – Aug 2023

- Architected **25+ Java Spring Boot** microservices exposing **REST APIs** for financial transactions for **1M+ users** at **99.9% uptime** on **AWS**.
- Designed high-throughput **SQL** pipelines on **AWS RDS** and event-driven messaging on **SQS**, cutting batch processing **8 min to 90 s** via **indexed views** and **partitioning**.
- Standardized **REST API** contracts and failure handling across **25+ services** (retry, timeout, propagation), reducing production escalations **60%** in one quarter.
- Mentored engineers** on **SDLC** practices and **code review** standards; wrote **Oracle PL/SQL** views and stored procedures for production **ETL**.

## PROJECTS

---

### LangFetch

[tinyurl.com/LangFetch](https://tinyurl.com/LangFetch)

- React/TypeScript** frontend, **Node.js** backend, and **PostgreSQL**; **LangGraph** agent decomposes questions into dependency-ordered SQL subqueries via **REST API**.

### Watch-your-LLM

[github.com/ganeshhgupta/watch-your-llm](https://github.com/ganeshhgupta/watch-your-llm)

- Full-stack LLM eval platform: **FastAPI** backend, **React/TypeScript** and **HTML5/CSS** frontend, **PostgreSQL**; rubric scoring and latency tracking.

### Distributed KV Store / VPC Control Plane Simulator

[github.com/ganeshhgupta/distributed-kv-store](https://github.com/ganeshhgupta/distributed-kv-store)

- Raft-consensus** KV store with **linearizable reads** and **log compaction**, plus a **VPC control plane** simulator.

## EDUCATION

---

University of Texas at Arlington | *M.S. Computer Science, GPA 4.0/4.0*

Aug 2023 – May 2025

MCKV Institute of Engineering | *B.Tech Computer Science*

Jul 2016 – Jul 2020

## SKILLS

---

**Languages:** JavaScript, TypeScript, Python, Java, SQL, HTML/CSS

**Full Stack:** React, Node.js, NestJS, Next.js, REST APIs, Microservices, JSON, JUnit, Mockito, DS&A

**Data & Messaging:** PostgreSQL, Redis, MongoDB, DynamoDB, NoSQL, Distributed Messaging (SQS, Kafka)

**Cloud & DevOps:** AWS (Lambda, S3, EC2, RDS, SQS), Azure, GCP, Docker, Kubernetes, GitHub Actions, CI/CD, Agile/Scrum

## CERTIFICATIONS

---

AWS Certified Developer Associate | GCP Professional MLOps Engineer