

Ganesh Gupta

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

EXPERIENCE

Oracle | *Software Engineer, AI Platform*

Jun 2025 – Present

- Built RESTful APIs using Python/FastAPI for a multi-agent AI platform, integrating LLM APIs (GPT, Claude, Llama, Gemini) into production workflows handling millions of daily transactions at 99%+ uptime.
- Engineered a pluggable distributed executor routing pipeline nodes across Docker containers and AWS Lambda, backed by an S3 artifact store eliminating redundant recomputation.
- Owned full-stack delivery of backend services and APIs (Python/FastAPI, React/TypeScript); replaced polling with SSE for live metrics streaming, cutting redundant API calls 90%.
- Designed and optimized agentic AI workflows across multiple LLMs, selecting and tuning models per use case for production quality, latency, and cost.
- Participated in design and code reviews enforcing style and testability standards, contributed to internal documentation, and triaged production issues affecting service reliability.

QureAI | *ML Infrastructure Engineer*

Mar 2026 – Present

- Cut p99 service latency 420ms to 140ms across 50K+ daily requests by migrating synchronous serving to async vLLM on GCP Cloud Run with autoscaling.
- Built PySpark data processing jobs for 10K+ large-scale feature dataset preparation feeding downstream training and evaluation.

Nomura Research Institute Financial Technologies | *Software Engineer*

Sep 2019 – Aug 2023

- Built and operated 25+ production microservices in Java, C++, and SQL-backed distributed systems for 1M+ users, sustaining 99.9% uptime over 3 years on AWS (EC2, RDS, SQS).
- Designed public-facing REST and gRPC APIs with CLI tooling, prioritizing backward compatibility and usability for partner integrations.
- Cut nightly financial data aggregation blocking downstream settlement from 8 minutes to 90 seconds by redesigning SQL with indexed views and batch partitioning.
- Standardized failure-handling contracts across 25+ financial services, reducing production escalations 60% in one quarter via consistent retry, timeout, and propagation semantics.

PROJECTS

Watch-your-LLM

github.com/ganeshhgupta/watch-your-llm

- Production LLM eval and observability platform with rubric-based scoring, drift detection, and latency tracking across model versions; FastAPI backend, React/TypeScript dashboard, PostgreSQL.

LangFetch

tinyurl.com/LangFetch

- LangGraph agent decomposing multi-table questions into dependency-ordered SQL subqueries with intermediate result passing across PostgreSQL; deployable as an MCP-compatible server for enterprise data queries.

RootCause AI

github.com/ganeshhgupta/rootcause-ai

- LangGraph multi-agent system automating production incident diagnosis via tool-use chains; routes alerts to specialized sub-agents and synthesizes root cause reports, reducing manual toil.

EDUCATION

University of Texas at Arlington | *M.S. Computer Science, GPA 4.0/4.0*

Aug 2023 – May 2025

MCKV Institute of Engineering | *B.Tech Computer Science*

Jul 2016 – Jul 2020

SKILLS

Languages: Python, Java, C++, TypeScript, SQL

Backend & APIs: FastAPI, REST/gRPC APIs, Microservices

AI/LLM: LLM APIs (GPT, Claude, Llama, Gemini), Agentic AI, Multi-Agent Systems (LangGraph), vLLM

Data & Big Data: PySpark, PostgreSQL, SQL/Relational Databases, Cassandra, DynamoDB, NoSQL

Cloud & DevOps: AWS (EC2, S3, Lambda), GCP (Cloud Run), Docker, Kubernetes, CI/CD

CERTIFICATIONS

AWS Certified Developer Associate | Google Cloud Professional Machine Learning Engineer