

Ganesh Gupta

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

EXPERIENCE

Oracle | *Software Engineer, AI Platform*

Jun 2025 – Present

- Architected agentic AI systems spanning multiple LLM providers (OpenAI, Anthropic, Google), driving model and framework selection via the OpenAI Agents SDK and MCP tool-use for production workloads at 99%+ uptime.
- Designed and led multi-agent orchestration patterns (LangGraph, CrewAI, AutoGen) across the engineering team, establishing planning, tool-calling, and failure-handling standards adopted org-wide.
- Drove the architecture for a RAG retrieval and serving pipeline on AWS SageMaker, deployed alongside Azure-hosted AI workloads, delivering real-time decisions at sub-200ms p99 latency.
- Owned the deployment-gated evaluation framework enforcing guardrails and quality checks across model outputs, blocking 15% of failed releases before production.

QureAI | *ML Infrastructure Engineer*

Mar 2026 – Present

- Led the redesign of a production RAG retrieval architecture, improving recall 61% to 89% on a 2K-query eval set via hybrid dense-sparse retrieval and re-ranking over a vector store.
- Drove multi-cloud GPU infrastructure strategy (AWS, GCP, CoreWeave) supporting distributed training and high-availability inference at production scale.

Nomura Research Institute Financial Technologies | *Software Engineer*

Sep 2019 – Aug 2023

- Architected 25+ production Java/C++ microservices for financial investment workflows serving 1M+ users at 99.9% uptime over 4 years on AWS, leading architecture decisions and code-review sessions.
- Standardized REST API failure-handling contracts across 25+ microservices, reducing escalations 60% via consistent retry, timeout, and authentication patterns.

PROJECTS

RootCause AI

github.com/ganeshhgupta/rootcause-ai

- Architected a LangGraph multi-agent system automating production incident diagnosis via tool-use chains; routes alerts to specialized sub-agents and synthesizes evidence-backed root cause reports.

LangFetch

tinyurl.com/LangFetch

- Designed a LangGraph agent with MCP-compatible tool-use decomposing natural language into dependency-ordered SQL subqueries; deployed as an API server for enterprise data queries.

Watch-your-LLM

github.com/ganeshhgupta/watch-your-llm

- Production LLM/agent evaluation platform with rubric-based scoring, drift detection, and regression alerting across model versions; FastAPI backend, PostgreSQL.

EDUCATION

University of Texas at Arlington | *M.S. Computer Science, GPA 4.0/4.0*

Aug 2023 – May 2025

MCKV Institute of Engineering | *B.Tech Computer Science*

Jul 2016 – Jul 2020

SKILLS

Languages: Python, Go, Java, C++, SQL

Agentic AI Architecture: LLMs/GenAI, MCP, Multi-Agent Systems (LangGraph, CrewAI, AutoGen), OpenAI Agents SDK, RAG, Prompt Engineering

Cloud: AWS (SageMaker, Lambda, EC2), Azure, GCP (Cloud Run, Vertex AI), Docker, Kubernetes

Production Systems: CI/CD, Evaluation/Guardrail Frameworks, Vector Databases (Pinecone), PostgreSQL

CERTIFICATIONS

AWS Certified Developer Associate | Google Cloud Professional Machine Learning Engineer