

Ganesh Gupta

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

EXPERIENCE

Oracle | *Software Engineer, AI Platform*

Jun 2025 – Present

- Designed and optimized agentic AI workflows across multiple LLMs (GPT, Claude, Llama, Gemini), selecting and tuning models per use case for production quality and cost.
- Built a multi-agent system via OpenAI Agents SDK and MCP tool-use handling millions of daily transactions at 99%+ uptime.
- Replaced polling with SSE for live ML metrics streaming, cutting redundant API calls 90% and enabling real-time model governance dashboards.
- Engineered a pluggable distributed executor routing pipeline nodes to Docker or AWS Lambda, backed by an S3 artifact store eliminating redundant recomputation.

QureAI | *ML Infrastructure Engineer*

Mar 2026 – Present

- Built a model observability framework tracking drift, latency, and eval regressions in production, enabling governance-aligned rollback workflows.
- Cut p99 inference latency 420ms to 140ms across 50K+ daily requests by migrating synchronous serving to async vLLM on GCP Cloud Run with autoscaling.

Nomura Research Institute Financial Technologies | *Software Engineer*

Sep 2019 – Aug 2023

- Built and operated 25+ production microservices in Java, C++, and SQL-backed distributed systems for 1M+ users, sustaining 99.9% uptime over 3 years on AWS (EC2, RDS, SQS).
- Cut nightly financial data aggregation blocking downstream settlement from 8 minutes to 90 seconds by redesigning SQL with indexed views and batch partitioning.
- Standardized failure-handling contracts across 25+ financial services, reducing production escalations 60% in one quarter via consistent retry, timeout, and propagation semantics.

PROJECTS

RootCause AI

github.com/ganeshhgupta/rootcause-ai

- Architected a LangGraph multi-agent system that autonomously diagnoses production incidents end-to-end, replacing manual log triage with policy-driven agentic investigation.
- Designed a tool-use chain enabling agents to pull logs, metrics, and traces on demand, then routes findings to specialized sub-agents scoped to distinct failure domains for parallel investigation.
- Built a synthesis layer that merges sub-agent findings into a single root-cause report with evidence trails, demonstrating agentic reasoning and orchestration directly applicable to enforcement-policy automation.

LLaMA 2 Fine-Tuning Pipeline for Financial Documents

github.com/ganeshhgupta/llama2-finance-finetuning

- Post-trained LLaMA 2 with LoRA and PEFT for domain-specific financial document understanding, reducing trainable parameter count while preserving model capacity via systematic hyperparameter optimization.

Watch-your-LLM

github.com/ganeshhgupta/watch-your-llm

- Production LLM eval platform with rubric-based scoring, drift detection, and latency tracking across model versions; FastAPI backend, React/TypeScript dashboard, PostgreSQL.

EDUCATION

University of Texas at Arlington | *M.S. Computer Science, GPA 4.0/4.0*

Aug 2023 – May 2025

MCKV Institute of Engineering | *B.Tech Computer Science*

Jul 2016 – Jul 2020

SKILLS

Languages: Python, Java, C++, TypeScript, SQL

AI/ML: LLMs (GPT, Claude, Llama, Gemini), LoRA/PEFT Fine-Tuning, Agentic AI, Multi-Agent Systems (LangGraph, OpenAI Agents SDK, MCP), Model Evaluation & Observability

Data: PostgreSQL, SQL/Relational Databases, NoSQL

Cloud & DevOps: AWS (EC2, S3, Lambda), GCP (Cloud Run), Docker, Kubernetes, CI/CD

CERTIFICATIONS

AWS Certified Developer Associate | Google Cloud Professional Machine Learning Engineer